

Generating Random Variates

Max Turgeon

STAT 3150–Statistical Computing

Lecture Objectives

- Recognize when to use the inverse-transform method.
- Be able to generate random variates through transformations.
- Derive bounding densities for accept-reject sampling.

Motivation

- A staple of modern statistical research is the **simulation study**.
 - Finite sample properties can then be compared to theoretical expectations.
- More generally, by simulating data we can study the properties of a method or a model.
- **Bayesian statistics** strongly relies on generating data to estimate the posterior density of the parameters (cf. STAT 4150).

- *Recall:* Let X be a random variable with CDF $F(x)$. The **quantile** function is defined as

$$F^{-1}(p) = \inf\{x \in \mathbb{R} \mid F(x) \geq p\}.$$

- If X is continuous, this is simply the inverse function.

Theorem

If U is uniform on $[0, 1]$, then $F^{-1}(U)$ has the same distribution as X .

- In R, we can sample random variates from $U(0, 1)$ by using the function `runif`:

```
runif(5)
```

```
## [1] 0.2681359 0.3308333 0.4411671 0.8352923  
0.9690489
```

Algorithm

To generate random variates from F :

1. Generate random variates from $U(0, 1)$.
2. Compute the quantile function F^{-1} .
3. Plug-in the uniform variates into F^{-1} .

Example i

- Let X follow an exponential distribution with parameter λ :

$$F(x) = 1 - \exp(-\lambda x).$$

- Since X is continuous, the quantile function is the inverse of F :

$$\begin{aligned} p = 1 - \exp(-\lambda x) &\Rightarrow \exp(-\lambda x) = 1 - p \\ &\Rightarrow -\lambda x = \log(1 - p) \\ &\Rightarrow x = \frac{-\log(1 - p)}{\lambda}. \end{aligned}$$


Example ii

```
lambda <- 1  
# We want 1000 samples  
n <- 1000  
unif_vars <- runif(1000)  
exp_vars <- -log(1 - unif_vars)/lambda
```


Example iii

```
# Compute theoretical quantiles
# using qexp
exp_theo <- qexp(ppoints(n))
qqplot(exp_theo, exp_vars)
# Add diagonal line
abline(a = 0, b = 1)
```

Example iv



Example v

Note: If U is uniform on $[0, 1]$, so is $1 - U$.

- Therefore $\frac{-\log(U)}{\lambda}$ also follows an $Exp(\lambda)$ distribution.

Inverse Transform—Discrete Edition





Figure 1: From Wikipedia

Example i

- Let X follow a Bernoulli distribution with parameter p :

$$F(x) = \begin{cases} 0 & x < 0, \\ 1 - p & x \in [0, 1), \\ 1 & x \geq 1. \end{cases}$$

Example ii



Example iii

- As we can see, we have

$$F^{-1}(u) = \begin{cases} 0 & u \leq 1 - p, \\ 1 & u > 1 - p. \end{cases}$$

- In other words, we sample U . If it is less than p , we set $X = 0$; else, we set $X = 1$.

Example iv

```
p <- 0.6
n <- 1000
unif_vars <- runif(1000)
# as.numeric turns FALSE into 0
# and TRUE into 1
bern_vars <- as.numeric(unif_vars > 1 - p)

c(mean(bern_vars), var(bern_vars))

## [1] 0.6400000 0.2306306
```


Example v

```
# Compare with theory
```

```
c(p, p*(1 - p))
```

```
## [1] 0.60 0.24
```

More General Transformations

- Inverse transform is just one type of transformation!
- We can use relationships between distributions to generate random variates. For example:
 - If $Z \sim N(0, 1)$, then $Z^2 \sim \chi^2(1)$.
 - If $V_1, \dots, V_p \sim \chi^2(1)$, then $\sum_{i=1}^p V_i \sim \chi^2(p)$.
 - If $U \sim \chi^2(p)$ and $V \sim \chi^2(q)$, then

$$\frac{U/p}{V/q} \sim F(p, q).$$

Example i

```
# Choose degrees of freedom
p <- 2
q <- 4

# rnorm samples from a normal distribution
U <- sum(rnorm(p)^2)
V <- sum(rnorm(q)^2)

# Take ratio
(U/p)/(V/q)
```

Example ii

```
## [1] 3.389027
```

```
# What if we want 1000 replicates?
```

```
# Use the function replicate!
```

```
# First argument: number of replicates
```


```
# Second argument: expression to be run multiple times
```

```
f_vars <- replicate(1000, {  
  U <- sum(rnorm(p)^2)  
  V <- sum(rnorm(q)^2)  
  (U/p)/(V/q)  
})
```

Example iii

```
qqplot(f_vars, qf(ppoints(1000), p, q))  
# Add diagonal line  
abline(a = 0, b = 1)
```

Example iv



Mixture of distributions i

- We say X with distribution function F follows a (discrete) **mixture of distributions** if we can write

$$F(x) = \sum_{k=1}^K \pi_k F_k(x),$$

where F_k is a distribution function for all k and $\sum_{k=1}^K \pi_k = 1$.

- For example, the distribution of adult heights can be well-approximated by a mixture of two normal distributions with equal weights.

Algorithm

To sample from $F = \sum_{k=1}^K \pi_k F_k$:

1. Select a component k by sampling from a multinomial distribution with probabilities (π_1, \dots, π_K) .
2. Sample from F_k .

Example i

- Assume that adult heights for **male** follows a normal distribution with mean 175cm and standard deviation 7cm.
- Assume that adult heights for **female** follows a normal distribution with mean 160cm and standard deviation 7cm.
- Finally assume that we have an equal proportion of male and female.
- How can we sample from the distribution of adult heights?

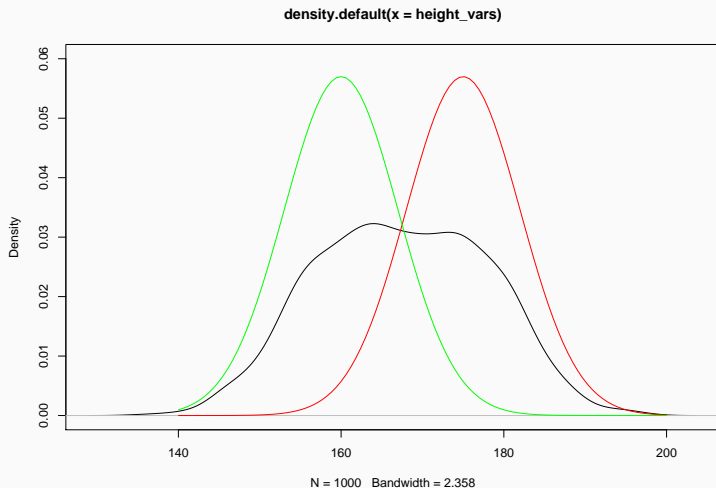
Example ii

```
height_vars <- replicate(1000, {  
  sex <- sample(c("M", "F"), size = 1,  
               prob = c(0.5, 0.5))  
  if (sex == "M") {  
    height <- rnorm(1, mean = 175, sd = 7)  
  }  
  if (sex == "F") {  
    height <- rnorm(1, mean = 160, sd = 7)  
  }  
  return(height)  
})
```

Example iii

```
plot(density(height_vars), ylim = c(0, 0.06))  
# Add component densities  
x_seq <- seq(140, 200, length.out = 100)  
lines(x_seq, dnorm(x_seq, 175, 7), col = "red")  
lines(x_seq, dnorm(x_seq, 160, 7), col = "green")
```

Example iv



Example v

```
# Sample with less code
sex <- sample(c(1, 2), size = 1000,
             replace = TRUE,
             prob = c(0.5, 0.5))
height_vars <- rnorm(1000, c(175, 160)[sex], 7)
c(mean(height_vars), sd(height_vars))

## [1] 167.8058 10.3375
```

Summary

- When we can compute the quantile function, the inverse transform is simple to implement.
 - But it can be hard to compute!
- We can leverage relationships between distributions to transform one random variate into another.
- **Mixture distributions** are a flexible way to model heterogeneous data, and it's “easy” to sample from them.
 - **Note:** Do not confuse *convolutions* (i.e. adding random variables) with *mixtures* (i.e. convex combination of distribution functions).

Acceptance-Reject Method i

- Suppose you want to sample from a distribution X with density f , but you can only sample from a different distribution Y with density g .
- Further suppose that there exists a constant $c > 1$ such that

$$\frac{f(t)}{g(t)} \leq c$$

for all t such that $f(t) > 0$.

- The *Acceptance-Reject method* is a way to transform random variates of Y into random variates of X .

Acceptance-Reject Method ii

Algorithm

1. Sample y from Y .
2. Sample a uniform variate u from $U(0, 1)$.
3. Compute the ratio $r := \frac{f(y)}{cg(y)}$. If $u < r$, set $x = y$. Otherwise, reject y and repeat from Step 1.

Note: The number of iterations before we accept a draw from Y follows a geometric distribution with mean c . So we want the constant c to be as small as possible.

(If you want a proof of why this works, see UM Learn.)


Example i

- We want to sample from $X \sim \text{Beta}(2, 2)$ whose density is $f(x) = 6x(1 - x)$.
 - The *proposal* distribution will be $Y \sim \text{Beta}(1, 1)$ (i.e. a uniform distribution).
- Let $t \in (0, 1)$. We have

$$\frac{f(t)}{g(t)} = \frac{6t(1 - t)}{1} \leq 6,$$

since the maximum t and $1 - t$ can take is 1. So we can set $c = 6$.

Example ii



Example iii

```
# Set parameters----  
C <- 6 # Constant  
n <- 1000 # Number of variates  
k <- 0 # counter for accepted  
j <- 0 # iterations  
y <- numeric(n) # Allocate memory
```

Example iv

```
# A while loop runs until condition no longer holds
while (k < n) {
  u <- runif(1)
  j <- j + 1
  x <- runif(1) # random variate from g
  if (u < 6*x*(1-x)/C) {
    k <- k + 1
    y[k] <- x
  }
}
```

Example v

```
# How many iterations did we need?  
j
```

```
## [1] 6144
```

```
# Compare theoretical and empirical quantiles  
p <- seq(0.1, 0.9, by = 0.1)  
Qhat <- quantile(y, p) # empirical  
Q <- qbeta(p, 2, 2) # theoretical
```

Example vi


```
round(cbind(Qhat, Q, diff = abs(Qhat - Q)), 3)
```

		Qhat	Q	diff
##	10%	0.181	0.196	0.014
##	20%	0.262	0.287	0.025
##	30%	0.341	0.363	0.022
##	40%	0.409	0.433	0.024
##	50%	0.477	0.500	0.023
##	60%	0.558	0.567	0.009
##	70%	0.637	0.637	0.000
##	80%	0.715	0.713	0.002
##	90%	0.793	0.804	0.011

Example vii

- As the graph showed, the “Rejection” region is very large.
 - In fact, it is unnecessarily large.
- With a little bit of calculus, we can show that the maximum value of $6x(1 - x)$ is 1.5.
 - In other words, we can set the constant $c = 1.5$.
 - This means that we can sample from X while rejecting 4 times *less* often.

Example viii



Example ix

```
C <- 1.5; k <- j <- 0 # Reset counters
while (k < n) {
  u <- runif(1)
  j <- j + 1
  x <- runif(1)
  if (u < 6*x*(1-x)/C) {
    k <- k + 1
    y[k] <- x
  }
}
```

Example x

```
# How many iterations did we need this time?  
j
```

```
## [1] 1517
```